

# Click Traffic Fraud Detection in Mobile Application Advertisements

Naiara de Almeida Pantuza

21/01/2020

Project developed during the DSA BigData Analytics with R and Microsoft Azure course: available on the kaggle platform

## DESCRIPTION:

Fraud risk is everywhere, but for companies that advertise online, click fraud can happen at an overwhelming volume, resulting in misleading click data and wasted money. Ad channels can drive up costs by simply clicking on the ad at a large scale. China is the largest mobile market in the world and therefore suffers from huge volumes of fraudulent traffic. TalkingData, China's largest independent big data service platform, covers over 70% of active mobile devices nationwide. They've built an IP blacklist and device blacklist. In this project, I am supposed to build an algorithm that predicts whether a user will download an app after clicking a mobile app ad.

```
# Setting work directory  
setwd("/home/naiara/Documentos/DataScience/FCD/BigDataRAzure/Project_Fraud_Detection")
```

All datasets were downloaded in kaggle. I chose to work with 10,000,000 lines to avoid processing and storage problems on my computer.

```
# Importing train and test datasets  
Azure <- FALSE  
  
if (Azure) {  
  train <- maml.mapInputPort(1)  
  test <- maml.mapInputPort(2)  
} else {  
  # Getting the first 10.000.000 rows of train dataset  
  train <- read.csv(unz("train.csv.zip", "train.csv"), nrows = 10000000)  
  # Getting test dataset  
  test <- read.csv("test.csv")  
}
```

## Data cleaning and processing

Here I proceeded with cleaning and data processing.

```
# Checking if there are missing values.  
# Only "attributed_time" has missing values, once the apps may not be downloaded.  
any(is.na(train[,7]))
```

```
## [1] FALSE
```

```
any(is.na(test))

## [1] FALSE

# Modifying variable types for analysis:
# Converting categorical variables to factor
cnames <- c('ip', 'app', 'device', 'os', 'channel')
for (var in cnames) {
  train[,var] <- as.factor(train[,var])
  test[,var] <- as.factor(test[,var])
}

# Converting target variable "is_attributed" to factor
train$is_attributed <- as.factor(train$is_attributed)

# Converting temporal variables to datetime format
train$click_time <- as.POSIXct(train$click_time, tz = Sys.timezone())
train$attributed_time <- dplyr::na_if(train$attributed_time, "")
train$attributed_time <- as.POSIXct(train$attributed_time, tz = Sys.timezone())
test$click_time <- as.POSIXct(test$click_time, tz = Sys.timezone())

# Converting "click_id" from test dataset to factor
test$click_id <- as.factor(test$click_id)
```

## Visual analysis and data exploration

Here I visualized training and test datasets and their summaries.

```
# Train dataset head
head(train)

##      ip app device os channel      click_time attributed_time
## 1  83230   3     1  13    379 2017-11-06 14:32:21          <NA>
## 2  17357   3     1  19    379 2017-11-06 14:33:34          <NA>
## 3  35810   3     1  13    379 2017-11-06 14:34:12          <NA>
## 4  45745  14     1  13    478 2017-11-06 14:34:52          <NA>
## 5 161007   3     1  13    379 2017-11-06 14:35:08          <NA>
## 6  18787   3     1  16    379 2017-11-06 14:36:26          <NA>
##      is_attributed
## 1                0
## 2                0
## 3                0
## 4                0
## 5                0
## 6                0

# Test dataset head
head(test)

##      click_id      ip app device os channel      click_time
## 1          0   5744   9     1  3    107 2017-11-10 04:00:00
## 2          1 119901   9     1  3    466 2017-11-10 04:00:00
## 3          2  72287  21     1  19    128 2017-11-10 04:00:00
## 4          3  78477  15     1  13    111 2017-11-10 04:00:00
```

```
## 5      4 123080 12      1 13      328 2017-11-10 04:00:00
## 6      5 110769 18      1 13      107 2017-11-10 04:00:00
```

```
# Train dataset summary
str(train)
```

```
## 'data.frame': 10000000 obs. of 8 variables:
## $ ip      : Factor w/ 68740 levels "9","10","19",...: 18073 3760 7775 9962 43910 4119 22306 2...
## $ app     : Factor w/ 332 levels "0","1","2","3",...: 4 4 4 15 4 4 4 4 61 ...
## $ device  : Factor w/ 940 levels "0","1","2","4",...: 2 2 2 2 2 2 2 2 2 ...
## $ os      : Factor w/ 292 levels "0","1","2","3",...: 14 20 14 14 14 17 24 20 14 23 ...
## $ channel : Factor w/ 170 levels "0","3","4","5",...: 116 116 116 158 116 116 116 116 116 149 ...
## $ click_time : POSIXct, format: "2017-11-06 14:32:21" "2017-11-06 14:33:34" ...
## $ attributed_time: POSIXct, format: NA NA ...
## $ is_attributed : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
```

```
summary(train)
```

```
##      ip      app      device      os
## 73516 : 51711 12      :1291185 1      :9381146 19      :2410148
## 73487 : 51215 2      :1202534 2      : 456617 13      :2199778
## 5314  : 35073 15      :1181585 3032   : 104393 17      : 531695
## 5348  : 35004 3      :1170412 0      :  46476 18      : 483602
## 53454 : 25381 9      : 966839 59      :  1618 22      : 365576
## 105560 : 23289 18      : 917820 40      :   462 10      : 285907
## (Other):9778327 (Other):3269625 (Other): 9288 (Other):3723294
##      channel      click_time      attributed_time
## 245      : 793105 Min.      :2017-11-06 14:32:21 Min.      :2017-11-06 16:00:47
## 134      : 630888 1st Qu.:2017-11-06 17:06:58 1st Qu.:2017-11-06 19:59:39
## 259      : 469845 Median :2017-11-06 20:27:57 Median :2017-11-06 23:34:47
## 477      : 412559 Mean   :2017-11-06 20:15:02 Mean   :2017-11-07 00:10:04
## 121      : 402226 3rd Qu.:2017-11-06 23:16:56 3rd Qu.:2017-11-07 02:40:27
## 107      : 388035 Max.   :2017-11-07 00:12:03 Max.   :2017-11-07 15:59:53
## (Other):6903342 NA's      :9981283
## is_attributed
## 0:9981283
## 1: 18717
##
##
##
##
##
##
```

```
# Teste dataset summary
summary(test)
```

```
##      click_id      ip      app      device
## 0      :      1 5348      : 182522 9      :2872176 1      :17360269
## 1      :      1 5314      : 162935 12      :2306083 2      : 1041975
## 2      :      1 73516      : 69089 3      :2201000 0      : 258152
## 3      :      1 73487      : 68866 2      :2060903 3      : 76117
## 4      :      1 53454      : 61503 18      :1923024 5      : 8279
## 5      :      1 114276      : 52649 15      :1079113 59      : 2775
## (Other):18790463 (Other):18192905 (Other):6348170 (Other): 42902
##      os      channel      click_time
## 19      :4334532 107      : 1214650 Min.      :2017-11-10 04:00:00
```

```
## 13      :3959515   265    : 778244   1st Qu.:2017-11-10 05:27:21
## 17      : 960531   232    : 684938   Median :2017-11-10 10:03:52
## 18      : 870068   477    : 683101   Mean   :2017-11-10 09:43:00
## 22      : 773184   178    : 582524   3rd Qu.:2017-11-10 13:34:07
## 8       : 520612   153    : 566046   Max.   :2017-11-10 15:00:00
## (Other):7372027   (Other):14280966
```

Including Plots

Check unique values for categorical variables from train dataset sample.

```
# Loading required packages
library(ggplot2)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

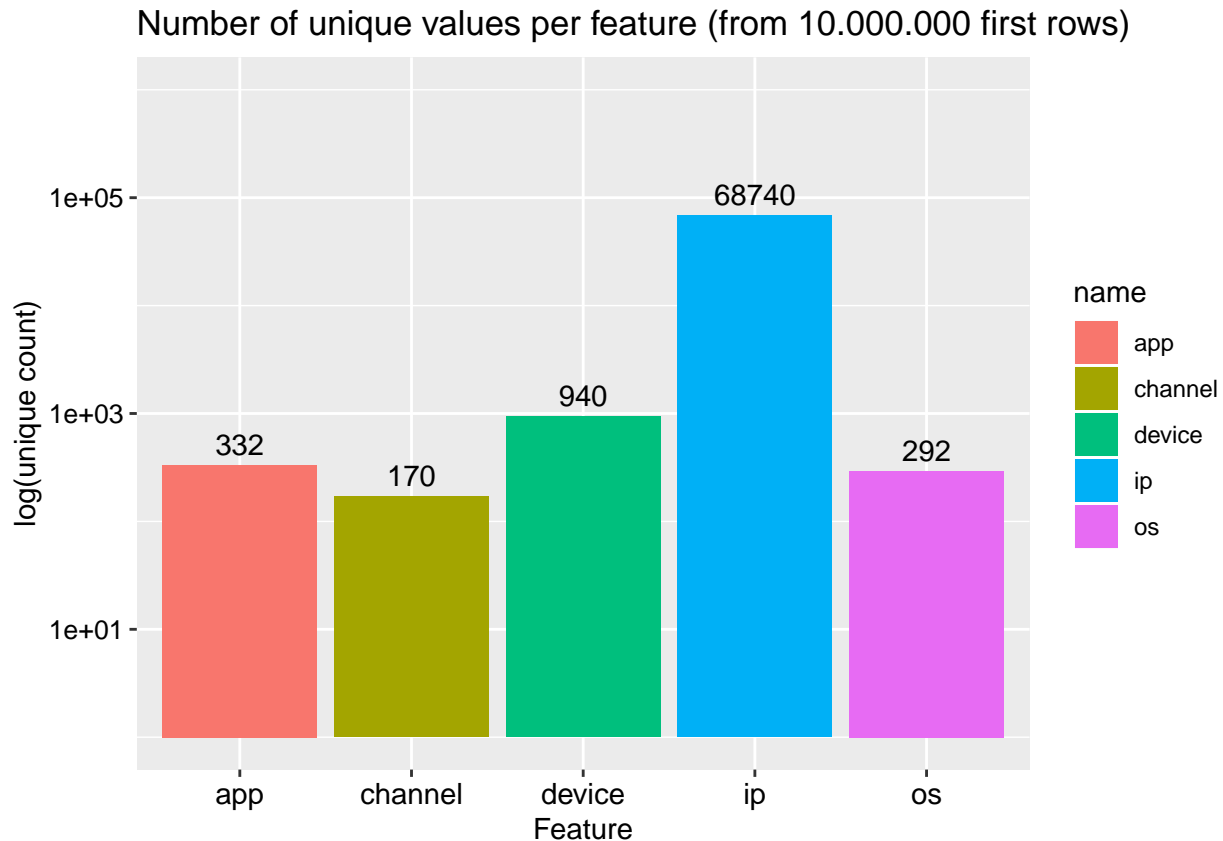
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

# Get the number of unique values for categoriacal dependent variables
uniq <- sapply(train[,cnames], function(x) {return(length(unique(x)))})
uniq_data <- data.frame(value=uniq, name=names(uniq), row.names = NULL)
```

Plotting the number of unique values for categoriacal dependent variables

```
# Bar plot of the number of single values
uniq_plot <- ggplot(uniq_data, aes(x=name, y=value, fill=name)) +
  geom_bar(stat = "identity") +
  ggtitle("Number of unique values per feature (from 10.000.000 first rows)") +
  xlab("Feature") +
  ylab("log(unique count)") +
  scale_y_log10(limits = c(1,1e6)) +
  geom_text(aes(label = sprintf("%d", value), y=value), vjust = -0.5) +
  theme(axis.text.x = element_text(hjust = 0.5, size=11,color="black")) +
  theme(axis.text.y = element_text(hjust = 0.5, size=10,color="black"))

uniq_plot
```



As noted in the graph, the variables have many levels.

Summarizing attribute variables from downloaded application logs.

```
# Filtering and summarizing attribute variables from downloaded application logs
downloaded <- train[train$is_attributed == "1", c("attributed_time", "is_attributed")]
summary(downloaded)
```

```
## attributed_time      is_attributed
## Min.   :2017-11-06 16:00:47 0:    0
## 1st Qu.:2017-11-06 19:59:39 1:18717
## Median :2017-11-06 23:34:47
## Mean   :2017-11-07 00:10:04
## 3rd Qu.:2017-11-07 02:40:27
## Max.   :2017-11-07 15:59:53
```


Plot proportion of apps download and not download

```
# Plotting proportion of apps download and not download
app_prop <- count(group_by(train[,c("app", "is_attributed")], is_attributed))
app_prop$n <- round(app_prop$n/sum(app_prop$n)*100, 2)
app_prop$is_attributed <- c("Not Downloaded (0)", "App Downloaded (1)")

appPropPlot <- ggplot(app_prop, aes(x=is_attributed, y=n)) +
  geom_bar(stat = "identity", fill="darkred") +
  ggtitle("App Downloaded x Not Downloades") +
  ylab("Percentage") +
  ylim(0,120) +
  geom_text(aes(label = sprintf("%.2f%%", n), y=n), vjust = -0.5) +
  theme(axis.text.x = element_text(hjust = 0.5, size=11,color="black")) +
```

```
theme(axis.text.y = element_text(hjust = 0.5, size=10,color="black"))
```

appPropPlot



As observed less than one percent of clicks were converted to downloads.

Analysing ips frequency

```
# Count 10 most clicked ips and their respective frequencies
temp <- as.data.frame(table(train$ip))
colnames(temp) <- c("ip", "count")
temp <- temp[order(temp$count, decreasing = TRUE),]
rownames(temp) <- 1:length(rownames(temp))
temp <- temp[1:10,]
temp
```

```
##      ip count
## 1  73516 51711
## 2  73487 51215
## 3   5314 35073
## 4   5348 35004
## 5  53454 25381
## 6 105560 23289
## 7 100275 23070
## 8 114276 22774
## 9 201182 22719
## 10 105475 22047
```

```

# Count 10 most clicked downloaded ips and their respective frequencies
temp_downloaded <- as.data.frame(table(train[train$is_attributed == "1",]$ip))
colnames(temp_downloaded) <- c("downloaded_ip", "count_downloads")
temp_downloaded <- temp_downloaded[order(temp_downloaded$count, decreasing = TRUE),]
rownames(temp_downloaded) <- 1:length(rownames(temp_downloaded))
temp_downloaded <- temp_downloaded[1:10,]
temp_downloaded

##      downloaded_ip count_downloads
## 1          73487          56
## 2          73516          54
## 3           5314          26
## 4         201182          25
## 5           5348          24
## 6         100275          23
## 7         105475          22
## 8         105560          16
## 9          44744          15
## 10        123994          14

# Checking coincide between 10 most clicked apss and the 10 most clicked downloaded apss
table(c(temp$ip,temp_downloaded$downloaded_ip))

```

```

##
## 1162 1173 9778 11715 16005 16012 21708 22868 22885 24864 26950 63053
##      2    2    1    1    2    2    2    2    2    1    1    2

```

Eight ips with the highest number of clicks were most downloaded.

Get summary of downloaded app dataset

```

# Get statistic data from downloaded ip dataset
# Minimum, maximum, average, median, quartiles.
summary(as.integer(train[train$is_attributed == "1", "ip"]))

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##         5  16336   32828   33468   50388   68737

```

```

# Count values
length(as.numeric(train[train$is_attributed == "1", "ip"]))

```

```
## [1] 18717
```

```

# Count unique values
length(unique(as.integer(train[train$is_attributed == "1", "ip"])))

```

```
## [1] 16112
```

Analysing Most Popular IPs Getting the 300 most clicked ips, respective count of clickes and conversion rates (clicks converted to download).

Exibing table results:

```

# Converting "is_attributed" to numeric temporarily to plotting
train$is_attributed <- as.numeric(train$is_attributed)-1

# Conversion Rates over Counts of 300 Most Popular IPs
# Exibing table results
count_rate <- train %>%

```

```
group_by(ip) %>%
summarise(click_count=n(), prop_downloaded = round(sum(is_attributed)/n(), 4)) %>%
arrange(desc(click_count)) %>%
slice(1:300) %>%
data.frame()
```

count\_rate

##	ip	click_count	prop_downloaded
## 1	73516	51711	0.0010
## 2	73487	51215	0.0011
## 3	5314	35073	0.0007
## 4	5348	35004	0.0007
## 5	53454	25381	0.0001
## 6	105560	23289	0.0007
## 7	100275	23070	0.0010
## 8	114276	22774	0.0001
## 9	201182	22719	0.0011
## 10	105475	22047	0.0010
## 11	95766	21966	0.0005
## 12	26995	19166	0.0005
## 13	209663	17605	0.0005
## 14	43793	15398	0.0008
## 15	137052	14840	0.0008
## 16	86767	14742	0.0004
## 17	17149	14673	0.0009
## 18	111025	14493	0.0001
## 19	138561	14119	0.0003
## 20	147957	14012	0.0010
## 21	114220	12818	0.0000
## 22	93054	12331	0.0009
## 23	92766	10904	0.0005
## 24	93021	10698	0.0010
## 25	92735	10534	0.0007
## 26	194308	9453	0.0006
## 27	93587	9450	0.0004
## 28	45745	9395	0.0011
## 29	44744	9232	0.0016
## 30	77048	9208	0.0014
## 31	114314	8831	0.0006
## 32	44725	8633	0.0012
## 33	84896	8594	0.0007
## 34	188387	8460	0.0011
## 35	114235	8431	0.0000
## 36	48240	8415	0.0002
## 37	194289	8239	0.0006
## 38	175837	8053	0.0004
## 39	48212	7995	0.0000
## 40	133522	7941	0.0013
## 41	3994	7884	0.0010
## 42	123994	7796	0.0018
## 43	48282	7774	0.0004
## 44	36150	7733	0.0004
## 45	4019	7719	0.0006



## 46	79881	7666	0.0005
## 47	43827	7652	0.0004
## 48	79857	7604	0.0005
## 49	178851	7425	0.0004
## 50	36183	7392	0.0004
## 51	44067	7383	0.0008
## 52	48170	7373	0.0000
## 53	147153	7319	0.0005
## 54	36213	7310	0.0010
## 55	108881	7265	0.0010
## 56	147046	7220	0.0008
## 57	125222	7184	0.0004
## 58	108913	7071	0.0006
## 59	147164	7034	0.0003
## 60	4052	6878	0.0009
## 61	147065	6849	0.0004
## 62	100393	6838	0.0010
## 63	105587	6701	0.0012
## 64	178873	6570	0.0002
## 65	5729	6557	0.0008
## 66	59125	6554	0.0002
## 67	79909	6541	0.0003
## 68	108341	6522	0.0005
## 69	108858	6456	0.0003
## 70	3964	6279	0.0008
## 71	79827	6207	0.0005
## 72	13634	6197	0.0008
## 73	119289	6177	0.0015
## 74	119369	5988	0.0008
## 75	90485	5981	0.0010
## 76	185670	5980	0.0005
## 77	108942	5975	0.0007
## 78	146001	5883	0.0003
## 79	52024	5814	0.0005
## 80	25071	5808	0.0005
## 81	52010	5704	0.0005
## 82	4989	5632	0.0018
## 83	37515	5411	0.0004
## 84	51992	5405	0.0009
## 85	109743	5375	0.0007
## 86	52043	5231	0.0008
## 87	84644	5195	0.0019
## 88	53715	5123	0.0004
## 89	84774	5104	0.0012
## 90	53964	5094	0.0004
## 91	119349	4984	0.0008
## 92	13597	4974	0.0004
## 93	100971	4959	0.0012
## 94	25097	4937	0.0014
## 95	197093	4871	0.0008
## 96	43855	4836	0.0010
## 97	90891	4742	0.0004
## 98	95820	4688	0.0006
## 99	90855	4678	0.0004

## 100	109723	4655	0.0002
## 101	85329	4613	0.0011
## 102	90509	4561	0.0009
## 103	25737	4451	0.0000
## 104	44673	4425	0.0014
## 105	44494	4383	0.0009
## 106	59395	4343	0.0005
## 107	92873	4331	0.0002
## 108	44458	4292	0.0009
## 109	97744	4206	0.0021
## 110	172483	4160	0.0012
## 111	133825	4157	0.0017
## 112	30587	4124	0.0010
## 113	114878	4090	0.0000
## 114	105910	4066	0.0000
## 115	172498	4066	0.0005
## 116	144604	4052	0.0000
## 117	85625	4035	0.0000
## 118	135992	4004	0.0000
## 119	91661	3944	0.0008
## 120	92673	3923	0.0003
## 121	91694	3910	0.0005
## 122	105433	3893	0.0005
## 123	105534	3890	0.0000
## 124	53960	3888	0.0003
## 125	87879	3869	0.0000
## 126	92852	3841	0.0005
## 127	70522	3822	0.0008
## 128	44555	3812	0.0010
## 129	151574	3805	0.0016
## 130	174548	3783	0.0019
## 131	75007	3776	0.0008
## 132	143418	3744	0.0000
## 133	25614	3732	0.0003
## 134	85644	3724	0.0008
## 135	105649	3704	0.0005
## 136	114678	3692	0.0011
## 137	97773	3690	0.0014
## 138	37948	3685	0.0000
## 139	25553	3682	0.0003
## 140	37919	3679	0.0014
## 141	105323	3677	0.0005
## 142	25679	3672	0.0003
## 143	87073	3659	0.0027
## 144	91712	3650	0.0003
## 145	97716	3637	0.0019
## 146	172522	3630	0.0006
## 147	105603	3629	0.0000
## 148	92712	3624	0.0011
## 149	105456	3616	0.0003
## 150	100182	3570	0.0006
## 151	192756	3551	0.0008
## 152	105569	3544	0.0011
## 153	91536	3521	0.0009

## 154	44527	3497	0.0014
## 155	172465	3493	0.0009
## 156	37972	3489	0.0009
## 157	25792	3468	0.0003
## 158	37892	3455	0.0009
## 159	37774	3453	0.0006
## 160	18703	3448	0.0000
## 161	91611	3446	0.0006
## 162	39756	3430	0.0009
## 163	105485	3416	0.0006
## 164	30564	3405	0.0003
## 165	25761	3391	0.0000
## 166	67197	3377	0.0006
## 167	151908	3364	0.0000
## 168	2095	3358	0.0015
## 169	53929	3355	0.0006
## 170	53479	3346	0.0003
## 171	114490	3340	0.0000
## 172	165085	3331	0.0003
## 173	76919	3307	0.0015
## 174	117867	3298	0.0012
## 175	25818	3291	0.0009
## 176	100212	3287	0.0003
## 177	97684	3284	0.0015
## 178	91574	3282	0.0006
## 179	67658	3259	0.0009
## 180	114461	3259	0.0009
## 181	2076	3256	0.0006
## 182	39782	3254	0.0003
## 183	105519	3242	0.0003
## 184	25705	3237	0.0003
## 185	109735	3233	0.0003
## 186	105292	3225	0.0000
## 187	202954	3185	0.0000
## 188	37813	3164	0.0003
## 189	114655	3137	0.0013
## 190	91885	3133	0.0013
## 191	85107	3125	0.0000
## 192	121472	3121	0.0019
## 193	205164	3121	0.0003
## 194	106460	3109	0.0003
## 195	54125	3089	0.0000
## 196	203048	3080	0.0000
## 197	25588	3073	0.0007
## 198	54157	3063	0.0003
## 199	176799	3059	0.0007
## 200	67628	3054	0.0003
## 201	25648	3032	0.0003
## 202	44536	2996	0.0003
## 203	118315	2987	0.0010
## 204	105861	2978	0.0007
## 205	24985	2967	0.0013
## 206	118339	2963	0.0007
## 207	176758	2961	0.0003

##	208	114904	2960	0.0003
##	209	91734	2955	0.0014
##	210	117898	2954	0.0014
##	211	12505	2950	0.0003
##	212	118229	2924	0.0007
##	213	118252	2914	0.0003
##	214	38219	2899	0.0021
##	215	118284	2863	0.0007
##	216	26814	2858	0.0014
##	217	38265	2850	0.0007
##	218	50169	2842	0.0014
##	219	62094	2837	0.0011
##	220	73671	2819	0.0004
##	221	4405	2809	0.0007
##	222	49383	2803	0.0018
##	223	48062	2800	0.0004
##	224	50512	2779	0.0022
##	225	118367	2766	0.0004
##	226	50482	2747	0.0022
##	227	76855	2746	0.0007
##	228	123788	2744	0.0000
##	229	50136	2733	0.0004
##	230	42139	2721	0.0015
##	231	100959	2697	0.0004
##	232	67439	2676	0.0004
##	233	111182	2672	0.0022
##	234	30614	2659	0.0008
##	235	38300	2635	0.0015
##	236	77085	2629	0.0011
##	237	99754	2626	0.0011
##	238	41232	2613	0.0011
##	239	106437	2607	0.0012
##	240	49431	2598	0.0023
##	241	99856	2595	0.0012
##	242	106598	2582	0.0008
##	243	49462	2573	0.0008
##	244	76885	2566	0.0004
##	245	109703	2555	0.0000
##	246	123703	2555	0.0008
##	247	99915	2552	0.0012
##	248	99769	2537	0.0004
##	249	99944	2533	0.0016
##	250	111153	2520	0.0024
##	251	37490	2518	0.0008
##	252	40289	2515	0.0008
##	253	44663	2509	0.0012
##	254	114816	2508	0.0004
##	255	100042	2507	0.0008
##	256	102467	2505	0.0008
##	257	100929	2504	0.0008
##	258	106200	2497	0.0020
##	259	145896	2493	0.0008
##	260	102235	2491	0.0012
##	261	12479	2485	0.0004

## 262	102264	2483	0.0012
## 263	67467	2480	0.0000
## 264	40077	2478	0.0020
## 265	137397	2474	0.0000
## 266	193448	2474	0.0016
## 267	44488	2460	0.0008
## 268	8208	2459	0.0024
## 269	74013	2458	0.0012
## 270	99897	2457	0.0000
## 271	106524	2457	0.0004
## 272	44615	2446	0.0000
## 273	52094	2445	0.0004
## 274	67037	2443	0.0000
## 275	73954	2442	0.0025
## 276	100088	2439	0.0004
## 277	49407	2437	0.0029
## 278	114795	2432	0.0012
## 279	44645	2416	0.0008
## 280	67606	2411	0.0000
## 281	193434	2398	0.0008
## 282	12524	2397	0.0004
## 283	191846	2395	0.0000
## 284	106223	2393	0.0021
## 285	8259	2389	0.0008
## 286	40372	2388	0.0004
## 287	62129	2385	0.0000
## 288	116425	2378	0.0004
## 289	193406	2378	0.0004
## 290	145934	2370	0.0000
## 291	8391	2363	0.0017
## 292	32453	2362	0.0000
## 293	44590	2359	0.0004
## 294	109644	2359	0.0000
## 295	40056	2353	0.0008
## 296	123759	2353	0.0004
## 297	106308	2347	0.0017
## 298	2564	2344	0.0000
## 299	50197	2336	0.0021
## 300	123654	2335	0.0004

Exibing plot results:

```
# Obtain the coefficient so that both y axes are on the same scale
count_rate$num <- seq(1:nrow(count_rate))
MAX <- max(count_rate$click_count)
mx <- max(count_rate$prop_downloaded)
coef <- mx/MAX


# Disabling scientific notation in R
options(scipen = 999)

# Plotting Conversion Rates over Counts of 300 Most Popular IPs
ggplot(count_rate, aes(x=num)) +
  geom_line(aes(x=num, y=click_count), color="darkred") +
  geom_line(aes(x=num, y=prop_downloaded/coef), color="darkgreen") +
```

```

ggtitle("Conversion Rates over Counts of 300 Most Popular IPs") +
scale_y_continuous(
  # Features of the first axis
  name = "Number of clicks",
  # Add a second axis and specify its features
  sec.axis = sec_axis(~.*coef, name="Proportion Downloaded")
) +
theme(
  axis.title.y = element_text(color = "darkred", size=13),
  axis.title.y.right = element_text(color = "darkgreen", size=13),
  plot.title = element_text(hjust = 0.5)
) +
xlab("")

```



According to the graph, the number of clicks and conversion rate for ips are not significantly correlated.

Repeating the previous analysis for applications, operating systems, devices and channels.

Getting the 100 most clicked apps, respective count of clickes and conversion rates (clicks converted to download) and exibing table and plot results:

```

# Conversion Rates over Counts of 100 Most Popular Apps
count_rate <- train %>%
  group_by(app) %>%
  summarise(click_count=n(), prop_downloaded = round(sum(is_attributed)/n(), 4)) %>%
  arrange(desc(click_count)) %>%
  slice(1:100) %>%
  data.frame()

count_rate

```

##	app	click_count	prop_downloaded
## 1	12	1291185	0.0001
## 2	2	1202534	0.0004
## 3	15	1181585	0.0003
## 4	3	1170412	0.0006
## 5	9	966839	0.0009
## 6	18	917820	0.0004
## 7	14	507491	0.0005
## 8	1	391508	0.0003
## 9	8	364361	0.0014
## 10	21	223823	0.0001
## 11	13	203332	0.0001
## 12	20	174792	0.0020
## 13	24	156247	0.0006
## 14	11	152367	0.0015
## 15	23	148119	0.0000
## 16	6	147356	0.0002
## 17	64	127923	0.0003
## 18	26	126630	0.0005
## 19	25	104855	0.0001
## 20	27	76417	0.0005
## 21	28	76050	0.0001
## 22	17	50956	0.0004
## 23	10	41224	0.0146
## 24	19	35586	0.1526
## 25	22	20734	0.0003
## 26	29	19773	0.0758
## 27	5	15570	0.0158
## 28	151	12388	0.0000
## 29	160	6990	0.0001
## 30	36	6736	0.0068
## 31	32	4762	0.0042
## 32	82	4343	0.0136
## 33	35	3986	0.7958
## 34	150	3476	0.0037
## 35	80	3384	0.0000
## 36	183	2974	0.0000
## 37	58	2531	0.0008
## 38	88	2462	0.0000
## 39	45	2276	0.0453
## 40	46	1967	0.0020
## 41	33	1847	0.0000
## 42	208	1756	0.0011
## 43	103	1639	0.0067
## 44	4	1567	0.0000
## 45	74	1521	0.0026
## 46	109	1471	0.0000
## 47	38	1463	0.0000
## 48	55	1394	0.0172
## 49	65	1359	0.0022
## 50	215	1235	0.0000
## 51	536	1233	0.0000
## 52	72	1165	0.4464
## 53	110	940	0.0021

## 54	68	863	0.0000
## 55	83	860	0.0744
## 56	39	795	0.1799
## 57	56	788	0.0000
## 58	107	731	0.1573
## 59	66	676	0.2278
## 60	60	671	0.0298
## 61	94	667	0.0000
## 62	95	654	0.0000
## 63	122	650	0.0185
## 64	315	614	0.0098
## 65	52	566	0.0477
## 66	265	555	0.0450
## 67	134	549	0.0000
## 68	181	549	0.0000
## 69	419	538	0.0000
## 70	202	519	0.0443
## 71	231	518	0.0000
## 72	170	495	0.0000
## 73	86	455	0.0066
## 74	37	440	0.1750
## 75	53	439	0.0000
## 76	121	404	0.2079
## 77	84	379	0.3325
## 78	118	356	0.0000
## 79	91	352	0.0000
## 80	119	347	0.0000
## 81	93	344	0.0000
## 82	59	320	0.0000
## 83	50	294	0.3401
## 84	85	291	0.0000
## 85	49	283	0.0000
## 86	232	276	0.0000
## 87	100	261	0.0000
## 88	79	252	0.3333
## 89	145	251	0.3307
## 90	108	249	0.3293
## 91	185	243	0.0000
## 92	137	235	0.0000
## 93	47	228	0.0000
## 94	16	226	0.2965
## 95	266	223	0.0000
## 96	115	220	0.4864
## 97	363	200	0.0000
## 98	78	197	0.2132
## 99	81	191	0.0000
## 100	76	189	0.0000


```

# Obtain the coefficient so that both y axes are on the same scale
count_rate$num <- seq(1:nrow(count_rate))
MAX <- max(count_rate$click_count)
mx <- max(count_rate$prop_downloaded)
coef <- mx/MAX

```



```
# Plotting Conversion Rates over Counts of 100 Most Popular Apps
ggplot(count_rate, aes(x=num)) +
  geom_line(aes(x=num, y=click_count), color="darkred") +
  geom_line(aes(x=num, y=prop_downloaded/coef), color="darkgreen") +
  ggtitle("Conversion Rates over Counts of 100 Most Popular Apps") +
  scale_y_continuous(
    # Features of the first axis
    name = "Number of clicks",
    # Add a second axis and specify its features
    sec.axis = sec_axis(~.*coef, name="Proportion Downloaded")
  ) +
  theme(
    axis.title.y = element_text(color = "darkred", size=13),
    axis.title.y.right = element_text(color = "darkgreen", size=13),
    plot.title = element_text(hjust = 0.5)
  ) +
  xlab("")
```



According to the graph, the number of clicks and conversion rate for apps are not significantly correlated.

Getting the 100 most clicked Operational Systems, respective count of clickes and conversion rates (clicks converted to download) and exibing table and plot results:

```
# Conversion Rates over Counts of Most Popular Operational Systems
count_rate <- train %>%
  group_by(os) %>%
  summarise(click_count=n(), prop_downloaded = round(sum(is_attributed)/n(), 4)) %>%
  arrange(desc(click_count)) %>%
  slice(1:100) %>%
  data.frame()
```

## count\_rate

##	os	click_count	prop_downloaded
## 1	19	2410148	0.0015
## 2	13	2199778	0.0013
## 3	17	531695	0.0012
## 4	18	483602	0.0011
## 5	22	365576	0.0017
## 6	10	285907	0.0010
## 7	8	279549	0.0010
## 8	6	242799	0.0026
## 9	9	239377	0.0007
## 10	25	232143	0.0012
## 11	15	230832	0.0009
## 12	20	223820	0.0009
## 13	16	166165	0.0013
## 14	37	151274	0.0005
## 15	3	147970	0.0009
## 16	14	134127	0.0017
## 17	41	126565	0.0009
## 18	1	113395	0.0030
## 19	607	107442	0.0009
## 20	12	107005	0.0008
## 21	27	94188	0.0005
## 22	35	93578	0.0008
## 23	23	86880	0.0006
## 24	32	84824	0.0005
## 25	53	80319	0.0004
## 26	28	75093	0.0029
## 27	11	70077	0.0004
## 28	47	63726	0.0002
## 29	30	58910	0.0004
## 30	26	47956	0.0010
## 31	31	38943	0.0013
## 32	2	38347	0.0002
## 33	36	36951	0.0003
## 34	49	35023	0.0013
## 35	40	32891	0.0102
## 36	4	30729	0.0006
## 37	42	21096	0.0000
## 38	0	17102	0.0947
## 39	43	16637	0.0002
## 40	34	15222	0.0020
## 41	58	14784	0.0000
## 42	46	14528	0.0003
## 43	24	13790	0.1437
## 44	7	11410	0.0004
## 45	38	9802	0.0478
## 46	48	9237	0.0011
## 47	44	8951	0.0009
## 48	5	8672	0.0000
## 49	65	7788	0.0000
## 50	56	7574	0.0000

## 51	70	6253	0.0000
## 52	79	5422	0.0000
## 53	21	4449	0.1875
## 54	66	4363	0.0002
## 55	64	3898	0.0000
## 56	29	3853	0.2027
## 57	52	3788	0.0000
## 58	55	3684	0.0000
## 59	50	3073	0.0843
## 60	77	3072	0.0000
## 61	39	3067	0.0000
## 62	73	3061	0.0000
## 63	97	2562	0.0000
## 64	62	1996	0.0000
## 65	63	1925	0.0000
## 66	76	1756	0.0154
## 67	98	1435	0.0000
## 68	90	1421	0.0000
## 69	59	1123	0.0436
## 70	57	1079	0.0000
## 71	96	1072	0.0028
## 72	109	1036	0.0019
## 73	100	913	0.0000
## 74	85	890	0.0000
## 75	60	865	0.0000
## 76	83	765	0.0000
## 77	74	475	0.0000
## 78	102	426	0.0000
## 79	69	407	0.0000
## 80	112	352	0.0000
## 81	71	339	0.0000
## 82	80	335	0.0000
## 83	84	294	0.0340
## 84	81	258	0.0078
## 85	87	241	0.0000
## 86	67	236	0.1992
## 87	106	236	0.0000
## 88	78	233	0.0000
## 89	92	225	0.0000
## 90	118	209	0.0000
## 91	111	192	0.0000
## 92	72	189	0.0000
## 93	107	171	0.0000
## 94	54	169	0.0000
## 95	68	165	0.0000
## 96	110	165	0.0000
## 97	155	154	0.0000
## 98	137	145	0.0000
## 99	89	142	0.0000
## 100	132	139	0.0000

```

# Obtain the coefficient so that both y axes are on the same scale
count_rate$num <- seq(1:nrow(count_rate))
MAX <- max(count_rate$click_count)


```

```

mx <- max(count_rate$prop_downloaded)
coef <- mx/MAX

# Plotting Conversion Rates over Counts of 100 Most Popular Operating Systems
ggplot(count_rate, aes(x=num)) +
  geom_line(aes(x=num, y=click_count), color="darkred") +
  geom_line(aes(x=num, y=prop_downloaded/coef), color="darkgreen") +
  ggtitle("Conversion Rates over Counts of 100 Most Popular Operating Systems") +
  scale_y_continuous(
    # Features of the first axis
    name = "Number of clicks",
    # Add a second axis and specify its features
    sec.axis = sec_axis(~.*coef, name="Proportion Downloaded")
  ) +
  theme(
    axis.title.y = element_text(color = "darkred", size=13),
    axis.title.y.right = element_text(color = "darkgreen", size=13),
    plot.title = element_text(hjust = 0.5)
  ) +
  xlab("")

```



According to the graph, the number of clicks and conversion rate for Operational Systems are not significantly correlated.

Getting the 100 most clicked devices, respective count of clickes and conversion rates (clicks converted to download) and exibing table and plot results:

```

# Conversion Rates and Counts of Most Popular by device
count_rate <- train %>%
  group_by(device) %>%

```

```

summarise(click_count=n(), prop_downloaded = round(sum(is_attributed)/n(), 4)) %>%
arrange(desc(click_count)) %>%
slice(1:100) %>%
data.frame()

```

count\_rate

##	device	click_count	prop_downloaded
## 1	1	9381146	0.0013
## 2	2	456617	0.0002
## 3	3032	104393	0.0000
## 4	0	46476	0.0920
## 5	59	1618	0.0012
## 6	40	462	0.2468
## 7	6	458	0.2227
## 8	16	334	0.2425
## 9	18	247	0.2267
## 10	33	204	0.1961
## 11	21	190	0.2421
## 12	154	151	0.1788
## 13	3033	151	0.1788
## 14	37	145	0.1931
## 15	30	126	0.3016
## 16	46	123	0.2114
## 17	114	122	0.1721
## 18	7	121	0.2314
## 19	88	117	0.3248
## 20	109	113	0.3009
## 21	67	111	0.2973
## 22	748	103	0.0000
## 23	136	96	0.2396
## 24	78	95	0.2000
## 25	82	95	0.2105
## 26	97	95	0.2211
## 27	374	95	0.0842
## 28	50	91	0.2747
## 29	211	81	0.2222
## 30	60	75	0.1867
## 31	203	75	0.1200
## 32	56	73	0.2192
## 33	96	72	0.1944
## 34	220	69	0.2029
## 35	343	69	0.1739
## 36	20	65	0.0615
## 37	4	60	0.2500
## 38	101	60	0.2500
## 39	214	58	0.1897
## 40	89	57	0.2456
## 41	231	51	0.1961
## 42	299	51	0.3137
## 43	73	50	0.1800
## 44	234	49	0.1633
## 45	168	48	0.1458
## 46	102	47	0.1064

## 47	103	47	0.2553
## 48	276	46	0.1522
## 49	137	45	0.2000
## 50	25	41	0.1707
## 51	127	41	0.1707
## 52	210	41	0.2439
## 53	11	40	0.2000
## 54	76	40	0.2250
## 55	36	37	0.1892
## 56	558	37	0.2432
## 57	42	36	0.1667
## 58	100	36	0.2778
## 59	189	36	0.1944
## 60	124	35	0.3143
## 61	95	34	0.3235
## 62	263	34	0.2059
## 63	381	34	0.3529
## 64	251	33	0.2121
## 65	736	33	0.0000
## 66	395	32	0.0625
## 67	9	31	0.1935
## 68	14	31	0.1613
## 69	75	31	0.1935
## 70	229	31	0.0968
## 71	350	31	0.2903
## 72	379	31	0.1935
## 73	479	30	0.0333
## 74	362	29	0.1379
## 75	52	28	0.0714
## 76	53	28	0.1786
## 77	54	27	0.1852
## 78	190	27	0.2222
## 79	51	26	0.1923
## 80	61	26	0.1154
## 81	208	26	0.1538
## 82	334	26	0.1154
## 83	68	25	0.2000
## 84	26	24	0.0417
## 85	396	24	0.2500
## 86	581	24	0.2083
## 87	118	23	0.2174
## 88	129	23	0.2609
## 89	160	23	0.1304
## 90	230	23	0.3913
## 91	338	23	0.2174
## 92	422	23	0.0870
## 93	19	22	0.1364
## 94	106	22	0.2273
## 95	169	22	0.1818
## 96	447	22	0.0909
## 97	486	22	0.1818
## 98	417	21	0.1905
## 99	132	20	0.3000
## 100	240	20	0.1500


```

# Obtain the coefficient so that both y axes are on the same scale
count_rate$num <- seq(1:nrow(count_rate))
MAX <- max(count_rate$click_count)
mx <- max(count_rate$prop_downloaded)
coef <- mx/MAX

# Plotting
ggplot(count_rate, aes(x=num)) +
  geom_line(aes(x=num, y=click_count), color="darkred") +
  geom_line(aes(x=num, y=prop_downloaded/coef), color="darkgreen") +
  ggtitle("Conversion Rates over Counts of 100 Most Popular Devices") +
  scale_y_continuous(
    # Features of the first axis
    name = "Number of clicks",
    # Add a second axis and specify its features
    sec.axis = sec_axis(~.*coef, name="Proportion Downloaded")
  ) +
  theme(
    axis.title.y = element_text(color = "darkred", size=13),
    axis.title.y.right = element_text(color = "darkgreen", size=13),
    plot.title = element_text(hjust = 0.5)
  ) +
  xlab("")

```

Conversion Rates over Counts of 100 Most Popular Devices



According to the graph, the number of clicks and conversion rate for devices are not significantly correlated.

Getting the 100 most clicked channels, respective count of clicks and conversion rates (clicks converted to download) and exhibiting table and plot results:

```
# Conversion Rates over Counts of Most Popular Channels
count_rate <- train %>%
  group_by(channel) %>%
  summarise(click_count=n(), prop_downloaded = round(sum(is_attributed)/n(), 4)) %>%
  arrange(desc(click_count)) %>%
  slice(1:100) %>%
  data.frame()

count_rate
```

##	channel	click_count	prop_downloaded
## 1	245	793105	0.0001
## 2	134	630888	0.0006
## 3	259	469845	0.0007
## 4	477	412559	0.0001
## 5	121	402226	0.0003
## 6	107	388035	0.0004
## 7	145	348862	0.0012
## 8	153	296832	0.0002
## 9	205	279720	0.0002
## 10	178	269720	0.0001
## 11	265	236949	0.0002
## 12	128	223205	0.0001
## 13	140	222096	0.0003
## 14	459	214060	0.0002
## 15	442	210687	0.0006
## 16	215	191618	0.0008
## 17	122	163312	0.0006
## 18	280	162425	0.0003
## 19	379	161608	0.0008
## 20	135	160215	0.0002
## 21	439	150074	0.0005
## 22	105	132516	0.0006
## 23	480	132030	0.0007
## 24	469	123869	0.0007
## 25	219	120673	0.0008
## 26	489	119261	0.0008
## 27	137	116107	0.0004
## 28	435	114372	0.0007
## 29	328	105392	0.0003
## 30	452	102074	0.0006
## 31	409	101462	0.0007
## 32	334	89689	0.0009
## 33	424	89682	0.0009
## 34	115	89512	0.0004
## 35	125	87895	0.0000
## 36	130	79525	0.0004
## 37	237	78336	0.0002
## 38	258	78247	0.0004
## 39	401	78163	0.0007
## 40	3	77703	0.0002
## 41	377	77054	0.0025
## 42	173	75319	0.0008
## 43	212	73234	0.0002




## 44	19	72148	0.0006
## 45	315	64963	0.0005
## 46	364	63361	0.0000
## 47	463	61366	0.0006
## 48	232	60869	0.0013
## 49	234	58134	0.0007
## 50	349	54608	0.0006
## 51	347	50608	0.0131
## 52	266	49840	0.0008
## 53	386	46399	0.0006
## 54	244	39003	0.0021
## 55	481	38936	0.0008
## 56	319	34687	0.0016
## 57	466	33805	0.0011
## 58	412	33307	0.0006
## 59	278	33270	0.0003
## 60	111	31583	0.0003
## 61	430	31557	0.0007
## 62	213	31500	0.1754
## 63	326	31000	0.0001
## 64	123	28072	0.0002
## 65	417	27184	0.0000
## 66	236	25486	0.0007
## 67	497	24556	0.0004
## 68	400	24432	0.0001
## 69	124	23929	0.0003
## 70	371	23466	0.0004
## 71	211	23380	0.0002
## 72	113	22058	0.0269
## 73	243	21376	0.0110
## 74	116	20162	0.0001
## 75	110	18514	0.0037
## 76	478	18278	0.0059
## 77	118	17917	0.0002
## 78	325	17344	0.0013
## 79	487	17344	0.0014
## 80	402	14470	0.0006
## 81	21	13911	0.0984
## 82	445	13104	0.0014
## 83	376	13037	0.0018
## 84	343	12989	0.0293
## 85	242	12197	0.0002
## 86	17	11958	0.0004
## 87	150	11685	0.0007
## 88	317	9752	0.0097
## 89	467	8287	0.0002
## 90	101	8143	0.1545
## 91	457	6987	0.0001
## 92	446	4077	0.0002
## 93	274	2703	0.7088
## 94	406	2615	0.0008
## 95	373	2514	0.0028
## 96	330	2482	0.0085
## 97	449	2352	0.0043

```
## 98      262      2208      0.0005
## 99      210      2200      0.0677
## 100     411      1829      0.0093
```

```
# Obtain the coefficient so that both y axes are on the same scale
count_rate$num <- seq(1:nrow(count_rate))
MAX <- max(count_rate$click_count)
mx <- max(count_rate$prop_downloaded)
coef <- mx/MAX

# Plotting
ggplot(count_rate, aes(x=num)) +
  geom_line(aes(x=num, y=click_count), color="darkred") +
  geom_line(aes(x=num, y=prop_downloaded/coef), color="darkgreen") +
  ggtitle("Conversion Rates over Counts of 100 Most Popular Channels") +
  scale_y_continuous(
    # Features of the first axis
    name = "Number of clicks",
    # Add a second axis and specify its features
    sec.axis = sec_axis(~.*coef, name="Proportion Downloaded")
  ) +
  theme(
    axis.title.y = element_text(color = "darkred", size=13),
    axis.title.y.right = element_text(color = "darkgreen", size=13),
    plot.title = element_text(hjust = 0.5)
  ) +
  xlab("")
```

Conversion Rates over Counts of 100 Most Popular Channels



According to the graph, the number of clicks and conversion rate for channels are not significantly correlated.

## Time Patterns

The analysis of temporal patterns will be made with the sample provided by the kaggle. The first lines of the dataset are organized by time and therefore are not random. Thus, they are inappropriate for detecting temporal patterns.

Getting and treating the random sample train dataset: converting datetime variables to POSIXct format and rounding "click\_time" variable's hours.

```
# Importing a random training dataset sample for time pattern analysis
sampleTrain <- read.csv("train_sample.csv")

head(sampleTrain)

##           ip app device os channel           click_time attributed_time
## 1  87540   12      1 13      497 2017-11-07 09:30:38
## 2 105560   25      1 17      259 2017-11-07 13:40:27
## 3 101424   12      1 19      212 2017-11-07 18:05:24
## 4  94584   13      1 13      477 2017-11-07 04:58:08
## 5  68413   12      1  1      178 2017-11-09 09:00:09
## 6  93663    3      1 17      115 2017-11-09 01:22:13
##    is_attributed
## 1                0
## 2                0
## 3                0
## 4                0
## 5                0
## 6                0

#convert click_time and attributed_time to time series
sampleTrain$attributed_time <- dplyr::na_if(sampleTrain$attributed_time, "")
sampleTrain$attributed_time <- as.POSIXct(sampleTrain$attributed_time, tz = Sys.timezone())
sampleTrain$click_time <- as.POSIXct(sampleTrain$click_time, tz = Sys.timezone())

# Convert "is_attributed" to numeric
sampleTrain$is_attributed <- as.numeric(sampleTrain$is_attributed)

#round the time to nearest hour
sampleTrain$click_round <- lubridate::round_date(sampleTrain$click_time, "hour")

# Checking for hourly patterns
count_rate <- sampleTrain[,c("click_round", "is_attributed")] %>%
  group_by(click_round) %>%
  summarise(click_count=n(), conversion_rate = round(sum(is_attributed)/n(), 4)) %>%
  data.frame()

head(count_rate, 100)

##           click_round click_count conversion_rate
## 1 2017-11-06 16:00:00         684          0.0000
## 2 2017-11-06 17:00:00         921          0.0022
## 3 2017-11-06 18:00:00         523          0.0000
## 4 2017-11-06 19:00:00         367          0.0000
## 5 2017-11-06 20:00:00         251          0.0000
## 6 2017-11-06 21:00:00         235          0.0085
```

## 7	2017-11-06 22:00:00	396	0.0000
## 8	2017-11-06 23:00:00	955	0.0021
## 9	2017-11-07 00:00:00	1649	0.0018
## 10	2017-11-07 01:00:00	1928	0.0031
## 11	2017-11-07 02:00:00	1715	0.0035
## 12	2017-11-07 03:00:00	1692	0.0053
## 13	2017-11-07 04:00:00	1861	0.0005
## 14	2017-11-07 05:00:00	1843	0.0016
## 15	2017-11-07 06:00:00	1700	0.0024
## 16	2017-11-07 07:00:00	1540	0.0019
## 17	2017-11-07 08:00:00	1609	0.0050
## 18	2017-11-07 09:00:00	1480	0.0014
## 19	2017-11-07 10:00:00	1645	0.0018
## 20	2017-11-07 11:00:00	1894	0.0042
## 21	2017-11-07 12:00:00	1631	0.0018
## 22	2017-11-07 13:00:00	1714	0.0023
## 23	2017-11-07 14:00:00	1763	0.0011
## 24	2017-11-07 15:00:00	1572	0.0019
## 25	2017-11-07 16:00:00	1462	0.0014
## 26	2017-11-07 17:00:00	939	0.0000
## 27	2017-11-07 18:00:00	514	0.0000
## 28	2017-11-07 19:00:00	297	0.0000
## 29	2017-11-07 20:00:00	242	0.0041
## 30	2017-11-07 21:00:00	218	0.0000
## 31	2017-11-07 22:00:00	417	0.0072
## 32	2017-11-07 23:00:00	1028	0.0049
## 33	2017-11-08 00:00:00	1692	0.0000
## 34	2017-11-08 01:00:00	1800	0.0044
## 35	2017-11-08 02:00:00	1688	0.0024
## 36	2017-11-08 03:00:00	1858	0.0032
## 37	2017-11-08 04:00:00	1788	0.0011
## 38	2017-11-08 05:00:00	1871	0.0043
## 39	2017-11-08 06:00:00	1607	0.0044
## 40	2017-11-08 07:00:00	1660	0.0024
## 41	2017-11-08 08:00:00	1667	0.0018
## 42	2017-11-08 09:00:00	1518	0.0007
## 43	2017-11-08 10:00:00	1884	0.0016
## 44	2017-11-08 11:00:00	1781	0.0034
## 45	2017-11-08 12:00:00	1937	0.0021
## 46	2017-11-08 13:00:00	1877	0.0016
## 47	2017-11-08 14:00:00	2030	0.0059
## 48	2017-11-08 15:00:00	1965	0.0005
## 49	2017-11-08 16:00:00	1564	0.0019
## 50	2017-11-08 17:00:00	1008	0.0000
## 51	2017-11-08 18:00:00	516	0.0019
## 52	2017-11-08 19:00:00	352	0.0028
## 53	2017-11-08 20:00:00	269	0.0037
## 54	2017-11-08 21:00:00	257	0.0039
## 55	2017-11-08 22:00:00	459	0.0000
## 56	2017-11-08 23:00:00	997	0.0000
## 57	2017-11-09 00:00:00	1599	0.0044
## 58	2017-11-09 01:00:00	1700	0.0018
## 59	2017-11-09 02:00:00	1609	0.0012
## 60	2017-11-09 03:00:00	1669	0.0000

```
## 61 2017-11-09 04:00:00      1963      0.0020
## 62 2017-11-09 05:00:00      2157      0.0023
## 63 2017-11-09 06:00:00      1962      0.0015
## 64 2017-11-09 07:00:00      1849      0.0022
## 65 2017-11-09 08:00:00      1571      0.0019
## 66 2017-11-09 09:00:00      1592      0.0057
## 67 2017-11-09 10:00:00      1684      0.0018
## 68 2017-11-09 11:00:00      1813      0.0022
## 69 2017-11-09 12:00:00      1748      0.0023
## 70 2017-11-09 13:00:00      1845      0.0027
## 71 2017-11-09 14:00:00      1943      0.0010
## 72 2017-11-09 15:00:00      1829      0.0027
## 73 2017-11-09 16:00:00       737      0.0000
```

```
# Plotting
```

```
require(gridExtra)
```

```
## Loading required package: gridExtra
```

```
##
```

```
## Attaching package: 'gridExtra'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      combine
```


```
# Plotting hourly clicks
```

```
plot1 <- ggplot(count_rate, aes(x=click_round)) +
  geom_line(aes(x=click_round, y=click_count), color="darkred") +
  ggtitle("Hourly click frequency") +
  scale_y_continuous(name = "Number of clicks") +
  scale_x_datetime(labels = scales::date_format("%H:%M\n%d-%b\n%Y"), date_breaks = "12 hours") +
  theme(
    axis.title.y = element_text(color = "black", size=13),
    axis.text.x= element_text(color = "black"),
    axis.text.y= element_text(color = "black"),
    plot.title = element_text(hjust = 0.5)
  )
```

```
# Plotting hourly conversion rate
```

```
plot2 <- ggplot(count_rate, aes(x=click_round)) +
  geom_line(aes(x=click_round, y=conversion_rate), color="darkblue") +
  ggtitle("Hourly conversion rate") +
  scale_y_continuous(name = "Conversion rate") +
  scale_x_datetime(labels = scales::date_format("%H:%M\n%d-%b\n%Y"), date_breaks = "12 hours",
    limits = ) +
  theme(
    axis.title.y = element_text(color = "black", size=13),
    axis.text.x= element_text(color = "black"),
    axis.text.y= element_text(color = "black"),
    plot.title = element_text(hjust = 0.5)
  )
```

```
grid.arrange(plot1, plot2, nrow = 2)
```



Looking at the graph, I noticed that the click count and conversion rate for downloads do not appear to be significantly correlated.

Extracting hour from "click\_time" variable and getting the number of clicks and conversion rate by hour. Exibing table and graphs:

```
# Extract hour from "click_time" and add to sample train dataset
sampleTrain$click_hour <- as.factor(lubridate::hour(sampleTrain$click_time))

# Getting number of clicks by hour
count_rate <- sampleTrain[,c("click_hour", "is_attributed")] %>%
  group_by(click_hour) %>%
  summarise(click_count=n(), conversion_rate = round(mean(is_attributed), 4)) %>%
  data.frame()


# Visualizing data.frame
View(count_rate)

# Plotting number of clicks by hour
ggplot(count_rate, aes(x=click_hour, y=click_count)) +
  geom_bar(stat = "identity", fill="purple", width = 0.5) +
  geom_line( aes(x=as.numeric(click_hour), y=click_count)) +
  ggtitle("Hourly click frequency") +
  scale_y_continuous(name = "Count of clicks") +
  theme(
    axis.title.y = element_text(color = "black", size=13),
    axis.text.x= element_text(color = "black"),
    axis.text.y= element_text(color = "black"),
```

```

plot.title = element_text(hjust = 0.5)
)


```



```

# Plotting hourly conversion rate
ggplot(count_rate, aes(x=click_hour, y=conversion_rate)) +
  geom_bar(stat = "identity", fill="purple", width = 0.5) +
  geom_line( aes(x=as.numeric(click_hour), y=conversion_rate)) +
  ggtitle("Hourly conversion rate") +
  scale_y_continuous(name = "Conversion rate") +
  theme(
    axis.title.y = element_text(color = "black", size=13),
    axis.text.x= element_text(color = "black"),
    axis.text.y= element_text(color = "black"),
    plot.title = element_text(hjust = 0.5)
  )
)

```




Apparently, fewer clicks occur between 27 and 22 hours and fewer downloads at 16, 17, 19 and 22 hours.

Plotting conversion rate and count of clicks to check if both variables correlate:

```
# Plotting conversion rate and count of clicks to check if the variables correlate
# Obtain the coefficient so that both y axes are on the same scale
MAX <- max(count_rate$click_count)
mx <- max(count_rate$conversion_rate)
coef <- mx/MAX

# Plotting
ggplot(count_rate, aes(x=num)) +
  geom_line(aes(x=as.numeric(click_hour), y=click_count, color="click time")) +
  geom_line(aes(x=as.numeric(click_hour), y=conversion_rate/coef, color="conversion rate")) +
  ggtitle("Conversion rates and click count by hour") +
  scale_y_continuous(
    # Features of the first axis
    name = "Count of clicks",
    # Add a second axis and specify its features
    sec.axis = sec_axis(~.*coef, name="Conversion Rate")
  ) + xlab("Hour") +
  theme(
    axis.title.y = element_text(color = "purple", size=13),
    axis.title.x = element_text(color = "black", size=13),
    axis.title.y.right = element_text(color = "darkred", size=13),
    plot.title = element_text(hjust = 0.5)
  ) +
  scale_color_manual(values = c("purple", "darkred"))
```





According to the graph, the variables appear to have a weak correlation.

Here, I check the time difference between click\_time and its conversion to download (attributed\_time) on sample train dataset: visualizing and summarizing

```
# Checking the time difference between clicking add and download it
sampleTrain$timeDiff <- hms::as_hms(sampleTrain$attributed_time - sampleTrain$click_time)

# Checking first rows and the time passed between click and download
head(sampleTrain[sampleTrain$is_attributed == 1,], 15)
```

##	ip	app	device	os	channel	click_time	attributed_time
## 285	224120	19	0	29	213	2017-11-08 02:22:13	2017-11-08 02:22:38
## 482	272894	10	1	7	113	2017-11-08 06:10:05	2017-11-08 06:10:37
## 1209	79001	19	0	0	213	2017-11-07 09:54:22	2017-11-07 11:59:05
## 1342	131029	19	0	0	343	2017-11-09 10:58:46	2017-11-09 11:52:01
## 1413	40352	19	0	0	213	2017-11-07 22:19:03	2017-11-08 01:55:02
## 1667	48733	35	1	18	274	2017-11-07 12:25:50	2017-11-07 13:10:30
## 1772	330861	35	1	22	21	2017-11-08 18:54:44	2017-11-08 22:39:52
## 1918	309576	5	1	32	113	2017-11-09 08:47:51	2017-11-09 08:47:55
## 3915	220571	71	1	25	3	2017-11-08 04:35:21	2017-11-08 04:37:46
## 3993	240051	35	1	19	21	2017-11-08 08:07:13	2017-11-08 09:46:42
## 4301	110652	19	16	0	213	2017-11-09 08:15:34	2017-11-09 09:30:19
## 4425	252612	5	1	31	113	2017-11-07 20:21:11	2017-11-07 20:21:42
## 4565	48072	19	21	24	213	2017-11-07 05:17:29	2017-11-07 06:49:01
## 4604	12506	62	1	19	21	2017-11-08 05:56:57	2017-11-08 08:56:58
## 4608	184467	35	1	30	274	2017-11-07 22:29:06	2017-11-08 00:16:14
##	is_attributed	click_round	click_hour	timeDiff			
## 285	1	2017-11-08	02:00:00	2	00:00:25		

```
## 482          1 2017-11-08 06:00:00          6 00:00:32
## 1209         1 2017-11-07 10:00:00          9 02:04:43
## 1342         1 2017-11-09 11:00:00         10 00:53:15
## 1413         1 2017-11-07 22:00:00        22 03:35:59
## 1667         1 2017-11-07 12:00:00        12 00:44:40
## 1772         1 2017-11-08 19:00:00        18 03:45:08
## 1918         1 2017-11-09 09:00:00         8 00:00:04
## 3915         1 2017-11-08 05:00:00         4 00:02:25
## 3993         1 2017-11-08 08:00:00         8 01:39:29
## 4301         1 2017-11-09 08:00:00         8 01:14:45
## 4425         1 2017-11-07 20:00:00        20 00:00:31
## 4565         1 2017-11-07 05:00:00         5 01:31:32
## 4604         1 2017-11-08 06:00:00         5 03:00:01
## 4608         1 2017-11-07 22:00:00        22 01:47:08
```

```
# Getting time passed summary
as.data.frame(lapply(summary(as.numeric(sampleTrain[sampleTrain$is_attributed == 1,
                                     "timeDiff"])), hms::as_hms))
```

```
##           Min.   X1st.Qu.   Median             Mean   X3rd.Qu.     Max.
## 1 00:00:02 00:00:52.5 00:03:18 01:14:59.572687 01:21:27.5 12:52:21
```

Here I check the time difference between click\_time and it's conversion to download (attributed\_time) on first 10.000.000 rows of train dataset:

```
# Checking the time difference between clicking add and download it
# on the first 10.000.000 rows of train dataset
train$timeDiff <- hms::as_hms(train$attributed_time - train$click_time)
```

```
# Summary
as.data.frame(lapply(summary(as.numeric(train[train$is_attributed == 1,
                                     "timeDiff"])), hms::as_hms))
```

```
##           Min. X1st.Qu.   Median             Mean X3rd.Qu.     Max.
## 1 00:00:00 00:01:26 00:25:03 03:48:02.893733 06:34:14 23:52:38
```

The difference varies from 0 to almost 24 hours (one day) on first rows of train dataset.

## Feature Selection

In this script the selection of characteristics is performed for the creation of the model. I used randomForest algorithm to measure variables' importance. So, I converted categorical dependent variables to integer. Due to my machine's processing and memory limitations, I worked with random samples from the training and test datasets to build and evaluate the model.

```
# Set seed
set.seed(123)

# Feature selection using randomForest package
# load the library
library(randomForest)
```

```
## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
```

```

## Attaching package: 'randomForest'

## The following object is masked from 'package:gridExtra':
##
##      combine

## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

# Converting target variable to factor
train$is_attributed <- as.factor(train$is_attributed)

# Converting dependents variables to integer to run random forest algorithm
cnames <- c('ip', 'app', 'device', 'os', 'channel')
for (var in cnames) {
  train[,var] <- as.integer(train[,var])
  test[,var] <- as.integer(test[,var])
}


# Due to my machine's processing and memory limitations, I will work with random samples
# from the training and test datasets to build and evaluate the model
index <- sample(1:nrow(train), 100000)
train_data <- train[index,]

# Creating random forest model measuring importance
model <- randomForest( is_attributed ~ ip + app + device + os + channel + click_time,
                      data = train_data,
                      ntree = 100,
                      nodesize = 10,
                      importance = TRUE)

# Plotting estimate variable importance
varImpPlot(model)

```

model



The dataset has few variables so I will make the first version of the model using all but (attribute\_time).

## Model construction and evaluation

In this script, the prediction model is created. Then, the prediction for the test data set is made and the model is evaluated. Due to my machine's processing and memory limitations, I will work with random samples from the training and test datasets on proportion 70% training to 30% test to build and evaluate the model. I used random forest algorithm again to build the model. Then I assessed its efficiency with a matrix of confusion, and accuracy. "sample\_submission.csv" file contains correct classification for test dataset.

Getting samples of train and test datasets and treating it.

```
# Getting the correct classification for test dataset
test_result <- read.csv("sample_submission.csv")

# Getting smaller samples of train and test datasets
# on proportion 70% training to 30% test
index <- sample(1:nrow(train), 100000)
train_data <- train[index, ]

str(train_data)
```

```
## 'data.frame': 100000 obs. of 9 variables:
## $ ip : int 1920 24864 3438 16667 9159 8045 43030 55274 15605 15727 ...
## $ app : int 13 14 9 3 4 7 19 13 13 3 ...
## $ device : int 2 2 2 2 2 2 2 2 2 2 ...
## $ os : int 19 7 14 7 7 18 20 14 14 21 ...
## $ channel : int 48 157 41 144 138 30 16 76 72 136 ...
## $ click_time : POSIXct, format: "2017-11-06 16:20:44" "2017-11-06 19:35:50" ...
```

```
## $ attributed_time: POSIXct, format: NA NA ...
## $ is_attributed : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ timeDiff      : 'hms' num NA NA NA NA ...
## ..- attr(*, "units")= chr "secs"

index <- sample(1:nrow(test), 42857)
test_data <- test[index,]
test_data <- merge(test_data, test_result[index,], by.x="click_id", by.y="click_id")
test_data$click_id <- NULL

# Converting dependent variables to integer before creating model
# using random forest algorithm
cnames <- c('ip', 'app', 'device', 'os', 'channel')
for (var in cnames) {
  train_data[,var] <- as.integer(train_data[,var])
  test_data[,var] <- as.integer(test_data[,var])
}

# Emphasize the levels of the target variable
# since I used a random sample to build the model
levels(train_data$is_attributed) <- c("0", "1")
```

Creating and printing model:

```
# Creating model
model <- randomForest( is_attributed ~ ip + app + device + os + channel + click_time,
  data = train_data,
  ntree = 100,
  nodesize = 10)

# Print model
print(model)
```

```
##
## Call:
## randomForest(formula = is_attributed ~ ip + app + device + os + channel + click_time, data = t
##           Type of random forest: classification
##           Number of trees: 100
## No. of variables tried at each split: 2
##
##           OOB estimate of error rate: 0.16%
## Confusion matrix:
##           0  1  class.error
## 0 99803  9 0.00009016952
## 1  154 34 0.81914893617
```

Predicting classification using sample test dataset and building a data.frame containing columns: expected results and correct results.

```
# Generating predictions in test data
pred <- data.frame(observed = test_data$is_attributed,
  predicted = predict(model, newdata = test_data[,1:6]))

pred$observed <- as.factor(pred$observed)
levels(pred$observed) = c("0", "1")
levels(pred$predicted) = c("0", "1")
```

```
# Visualizing the results  
View(pred)
```

Evaluating model

```
# Evaluating the model  
# Generating a confusion matrix  
caret::confusionMatrix(pred$observed, pred$predicted)
```

```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction      0      1  
##           0 42835     22  
##           1      0      0  
##  
##           Accuracy : 0.9995  
##           95% CI : (0.9992, 0.9997)  
##      No Information Rate : 0.9995  
##      P-Value [Acc > NIR] : 0.5564  
##  
##           Kappa : 0  
##  
##  McNemar's Test P-Value : 0.000007562  
##  
##           Sensitivity : 1.0000  
##           Specificity : 0.0000  
##      Pos Pred Value : 0.9995  
##      Neg Pred Value :      NaN  
##           Prevalence : 0.9995  
##      Detection Rate : 0.9995  
##      Detection Prevalence : 1.0000  
##      Balanced Accuracy : 0.5000  
##  
##      'Positive' Class : 0  
##
```

The model created has the model created has an accuracy of 0.9996. Besides, it's a efficient model. However as the tests were done with a small sample I cannot say that it is an efficient model. The sample may not contemplate positive oficial results (downloaded apps).